

Session-Aware Task Assistant with Long-Term Memory

Ajibade Ayomide <a.ajibade@innopolis.university>
Ali Salloum <a.salloum@innopolis.university>

Innopolis University
Agentic AI — Final Project Report

26 July 2026

Abstract

We present a ReAct-style task assistant that manages interaction-derived knowledge with dual-tier memory: a bounded short-term working context (scratchpad, rolling summary, selective turn recall) and a Chroma-backed long-term store with explicit **remember/recall** tools and post-task reflection. The agent is orchestrated with LangGraph and a SQLite checkpoint, instrumented with local JSON traces and optional Langfuse, and evaluated against two in-repo baselines (stateless and full-history). We report *live* OpenRouter metrics for task quality and an *offline harness* for checker sanity and context efficiency ($\approx 41\%$ token reduction vs. full history on a shared scenario core). Negative controls confirm that the evaluation checker rejects known failure modes.

Contents

1	Introduction	2
2	Problem and goals	2
3	Method	2
3.1	Agent loop	2
3.2	Short-term memory	2
3.3	Long-term memory	3
3.4	Orchestration and observability	3
3.5	Baselines	3
4	Architecture	3
5	Memory write policy	3
6	Evaluation	3
6.1	Offline harness (regression / checker)	4
6.2	Live system metrics (OpenRouter)	4
7	Discussion and limitations	4
8	Task distribution	4
9	Conclusion	4

1 Introduction

Large language model agents are effectively stateless: each call sees only the messages supplied in the current context window. Multi-turn and cross-session work therefore loses user preferences, constraints, and prior outcomes unless the user repeats them or the system retains an unbounded transcript. Static retrieval over documents grounds answers in a corpus but not in *interaction-derived* knowledge (for example, “prefer metric units” or “we rejected design A last week”).

This project implements a lightweight assistant that executes multi-step tasks with tools while managing memory explicitly: what to write, what to recall, and how to stay within a token budget—without relying on million-token windows. The design follows the course framing of profile, memory, planning, and action; Thought–Action–Observation loops with a harness-owned observation channel; and context treated as a budget rather than a dump.

Course deliverables. The submission consists of three artefacts:

1. **Report** — this LaTeX document (PDF).
2. **Repository** — source code at <https://github.com/ali-salloum6/Session-Aware-Task-Assistant>.
3. **Video** — a live demo of the project running on a real OpenRouter run (preference write/recall, tools, local traces, Langfuse, and evaluation tables).

2 Problem and goals

Goals.

1. Execute multi-step tasks with tools (`calculator`, `scoped search`, `remember`, `recall`, `update_scratchpad`).
2. Keep a **bounded** short-term context (scratchpad + summary + selective recall + last- N turns).
3. Persist durable facts across sessions in a vector store, scoped by `user_id`.
4. Compare against (B0) stateless ReAct and (B1) append-all history using a fixed scenario suite.
5. Produce reproducible artefacts: repository, LaTeX report, and a live demo video of a real run.

Non-goals. Production authentication, consent productization, multi-agent orchestration, GraphRAG, MCP tool hosting, and fine-tuning are out of scope.

3 Method

3.1 Agent loop

The product path uses native tool calling. Each step the model may emit tool calls; the harness executes tools and appends observations. The loop stops on a final answer or a hard step budget (`MAX_STEPS`). The default LLM is OpenRouter `google/gemini-2.5-flash`.

3.2 Short-term memory

Short-term memory (STM) maintains dialogue turns and a structured scratchpad (`PLAN` / `DONE` / `NEXT`) reinjected every model call. In **bounded** mode, when the estimated token budget is approached, older turns are compressed into a rolling summary and selectively recalled by relevance; recent turns remain verbatim. In **full** mode (baseline B1), history is appended without summarization.

3.3 Long-term memory

Long-term memory (LTM) is a Chroma collection on disk. Writes occur via the **remember** tool (explicit user request) or a single post-task reflection pass that may propose 0–3 durable facts and must abstain when nothing lasting was learned. Near-duplicates are skipped by cosine distance. Entries carry **memory_type** (episodic / semantic / procedural), source, and timestamps. Recall is top- k and user-scoped; empty recall must surface as unknown (no confabulation).

3.4 Orchestration and observability

LangGraph nodes implement: assemble context $\rightarrow$ agent $\rightarrow$ tools (loop) $\rightarrow$ finalize $\rightarrow$ reflect memory. Session resume uses a SQLite checkpointer keyed by **session_id**. Every run writes a local JSON trace (input, tools, recalled IDs, token estimates, latency). When Langfuse keys are configured, the same payload is emitted as an agent observation with nested tool spans.

3.5 Baselines

Mode	Alias	Behaviour
stateless	B0	No STM history, no LTM tools/injection, no reflection
full_history	B1	STM append-all, no LTM
memory	B2	Bounded STM + LTM + reflection

4 Architecture

Figure 1 summarises the LangGraph control flow used in memory mode.

START $\rightarrow$ assemble_context $\rightarrow$ agent $\rightleftharpoons$ tools $\rightarrow$ finalize $\rightarrow$ reflect_memory $\rightarrow$ END

Figure 1: LangGraph node pipeline (memory mode). Checkpointer: SQLite; thread_id=session_id.

Context assembly concatenates system instructions, scratchpad, optional LTM hits, STM history (summary + recalled turns + recent turns), and the current user task. Reflection runs only in memory mode and applies a deduplication gate before writes. Full notes live in the repository at docs/architecture.md and docs/memory-policy.md.

5 Memory write policy

Long-term memory is written deliberately: on explicit user request via **remember**, or after a task via reflection that may propose 0–3 durable facts (preferences, constraints, decisions) and must abstain when nothing lasting was learned. Near-duplicates are skipped so repeated preferences do not multiply. Ephemeral arithmetic and tool noise are not stored. Empty recall must be answered as unknown, never confabulated.

6 Evaluation

We separate two result layers and do not mix their percentages.

6.1 Offline harness (regression / checker)

Scripted tool calls with observation-only final answers exercise the metric checkers. Shared-core rows use equal N across B0/B1/B2. Negative controls inject missing tools, wrong search results, skipped remembers, and confabulation; the checker is expected to fail them.

Table 1: Offline harness (shared-core). Not model-quality scores.

Mode	N	Task compl.	Recall	Faithfulness	Mean tokens
stateless	10	100.0%	100.0%	100.0%	1563
full_history	10	100.0%	100.0%	100.0%	7033
memory	10	100.0%	100.0%	100.0%	4530

Context efficiency vs. B1 on the shared core: **41.1%** mean-token reduction for memory mode (target $\geq 30\%$). Negative controls: **4/4** expected failures correctly caught.

6.2 Live system metrics (OpenRouter)

A frozen live subset runs the real model end-to-end. These numbers support quality claims.

Table 2: Live system metrics (google/gemini-2.5-flash via OpenRouter).

Mode	N	Task compl.	Recall	Faithfulness	Mean tokens
stateless	4	100.0%	100.0%	100.0%	988
full_history	4	75.0%	100.0%	100.0%	1098
memory	8	100.0%	100.0%	100.0%	1162

Reading the tables. Cite Table 2 for task completion / recall / faithfulness. Cite Table 1 for context efficiency and checker sanity. The live full-history STM miss (75% task completion) is a real model failure retained in the report.

7 Discussion and limitations

- Live N is small (frozen report subset) by design for cost and deadline; harness N is larger but scripted.
- Token estimates are heuristic; absolute token counts from the provider may differ.
- Embeddings and Chroma packaging can vary across platforms; the store is local-directory persistent.
- Reflection quality depends on the LLM; policy and dedup mitigate, but do not eliminate, weak writes.
- Langfuse is optional; local JSON traces are the always-on ground truth for debugging.

8 Task distribution

9 Conclusion

We delivered a session-aware ReAct assistant with bounded short-term context, cross-session long-term memory, LangGraph orchestration, dual-layer evaluation, and reproducible run traces

Table 3: Phase ownership (lead / support).

Phase	Lead	Support
P0–P1 (smoke, ReAct tools)	Ajibade Ayomide	Ali Salloum
P2–P3 (scratchpad, context budget)	Ali Salloum	Ajibade Ayomide
P4–P5 (LTM, write policy)	Ajibade Ayomide	Ali Salloum
P6–P7 (LangGraph, observability)	Ali Salloum	Ajibade Ayomide
P8 scenarios	both (split categories)	—
P8 harness / metrics	Ajibade Ayomide	—
P9 report (draft)	Ali Salloum	—
P9 report (methods/eval)	Ajibade Ayomide	—
P9 video (stable live run)	Ali Salloum	—

(local and Langfuse). Live memory-mode runs meet the proposal quality targets on the frozen suite; the offline harness demonstrates context efficiency against full history and validates that the metric checker detects intentional failures.

Submission artefacts

The three course deliverables are:

1. **Report:** this PDF (`report/report.pdf`; source `report/report.tex`).
2. **Repository:** <https://github.com/ali-salloum6/Session-Aware-Task-Assistant> (includes `src/assistant/`, `scripts/`, `eval/`, `docs/architecture.md`, `docs/memory-policy.md`).
3. **Video:** live demo of a real run (bundled with the submission package).

References

- [1] S. Yao et al., “ReAct: Synergizing Reasoning and Acting in Language Models,” ICLR, 2023.
- [2] LangChain, “LangGraph,” documentation, 2024–2026.
- [3] Langfuse, “Open-source LLM engineering platform,” documentation, 2024–2026.
- [4] Chroma, “Chroma vector database,” documentation, 2024–2026.